

What to measure?

- *Stress testing* or *load testing*: how far can I push my system...
- ...before performance becomes unacceptable?
- ...before it gasps and dies?
- Usually, one component will be *bottleneck*
- a particular view, action, query, ...
- Load testers can be simple or sophisticated
- bang on a single URI over and over
- do a fixed sequence of URI's over and over
- play back a log file

Longevity Bugs

- Resource leak (RAM, file buffers, sessions table) is classic example
 - Some infrastructure software such as Apache already does *rejuvenation*
 - aka “rolling reboot”
 - Related: running out of sessions
 - Solution: store whole `session[]` in cookie (Rails 3 does this by default)
-

Which is probably *not* a metric of high interest to you, the app operator?

- Slowest queries
- Maximum CPU utilization
- 99%ile response time
- Rendering time of 3 slowest views



The fastest database is the one you don't use

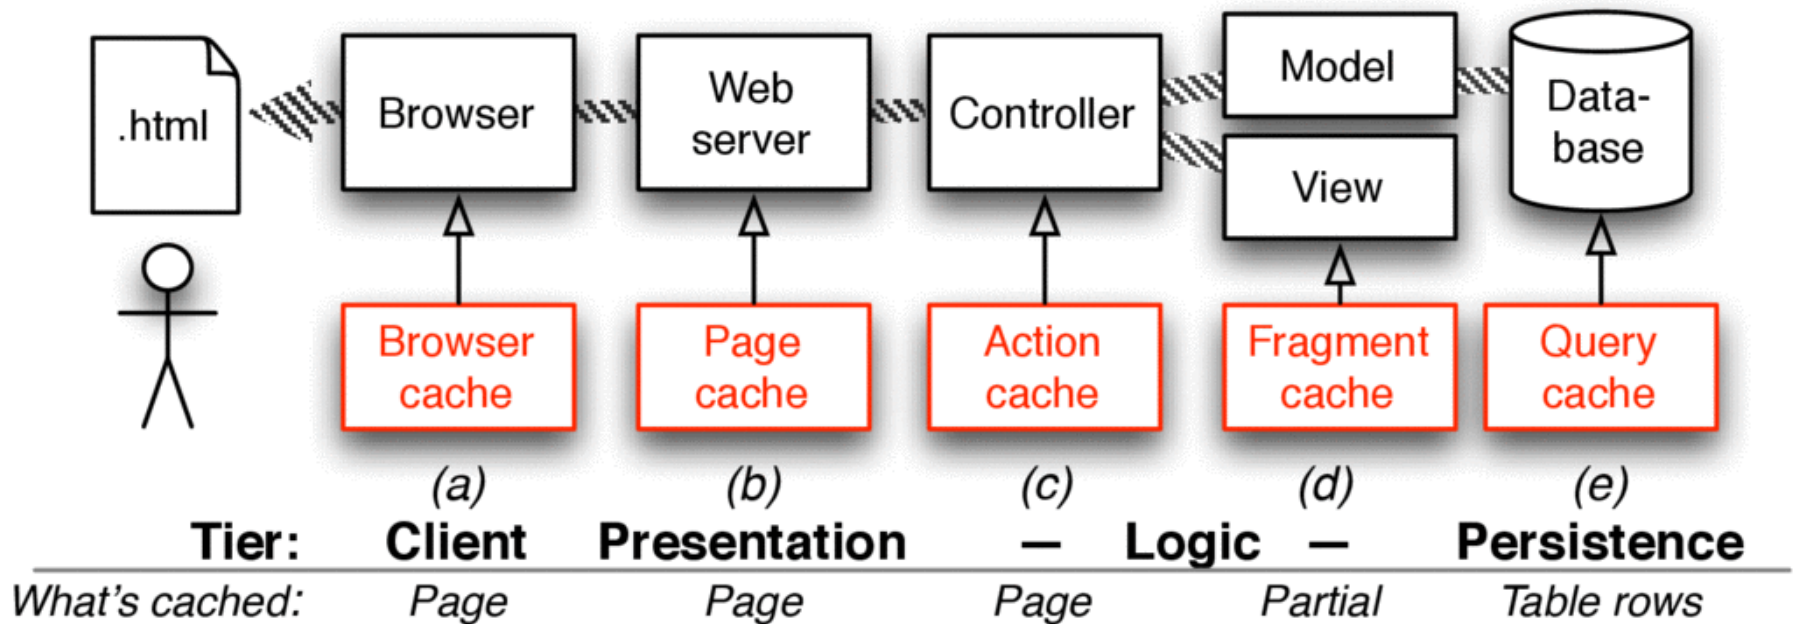
- **Caching:** Avoid touching database if answer to a query hasn't changed

1. Identify what to cache

- whole view: page & action caching
- parts of view: fragment caching with partials

2. Invalidate (get rid of) *stale* cached versions when underlying DB changes

Cache flow



Page & Action Caching

- When: output of *entire action* can be cached
- *Page caching* bypasses controller action
- *Action caching* runs filters first
- *Corollary*: don't mix filter & non-filter code paths in same action!
- *Caveat*: caching based on page URL *without* optional "?" parameters!

/movies/index?rating=PG = movies/index

/movies/index/rating/PG ≠ movies/index

Example

- Bad:

```
cache_page :index
def index
  if logged_in?
    ...
  else
    redirect_to login_path
  end
end
```

- Better:

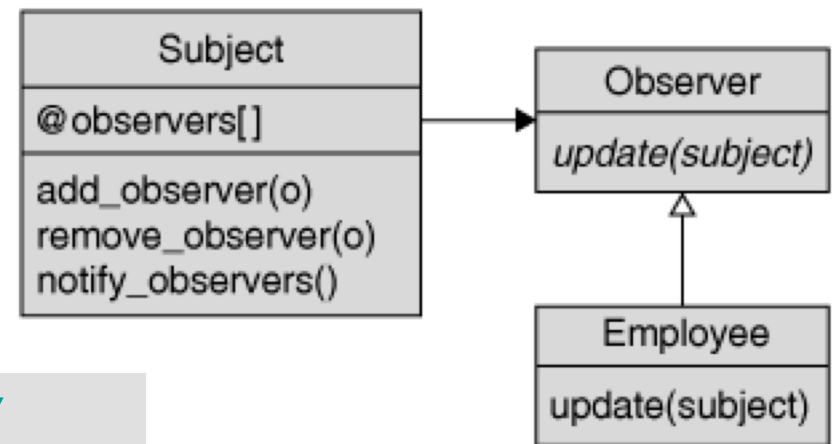
```
cache_page :public_index
cache_action :logged_in_index
before_filter :check_logged_in,
:only => 'logged_in_index'
def public_index
  ...
end

def logged_in_index
  ...
end
```

Fragment caching for views

- Caches HTML resulting from rendering part of a page (e.g. partial)
- Cuts back on *rendering time*, not database queries
- cache "movies_with_ratings" do
= render :collection => @movies


- How do we detect when cached versions no longer match database?
- Sweepers use Observer design pattern to separate expiration logic from rest of app



<http://pastebin.com/fCZJSimS>

How much does caching help?

- With ~1K movies and ~100 reviews/movie in RottenPotatoes on Heroku, **heroku logs** shows:



- Can serve 8x to 21x as many users with same number of servers if caching used

Under-17 visitors to RottenPotatoes shouldn't see NC-17 movies in any listing. What kinds of caching would be appropriate:

i) Page ii) Action iii) Fragment

- ☐ (i) & (iii)
- ☐ (ii) & (iii)
- ☐ (iii) only
- ☐ (i), (ii) and (iii)



Be kind to the database

- Outgrowing a single-machine database means a big investment: sharding, replication, etc.
 - Alternative: find ways to relieve pressure on database so can stay in “PaaS-friendly” tier
 1. Use **caching** to reduce number of database accesses
 2. Avoid “**n+1 queries**” **problem** in Associations
 3. Use **indices** judiciously
-

n+1 queries problem

- **Problem:** you are doing $n+1$ queries to traverse an association, rather than 1 query

<http://pastebin.com/QKxqcbhk>

- **Solution:** bullet gem can help you find these



- **Lesson:** all abstractions eventually leak!
-

Eager loading

- Naive way:

```
@movie = movie.where( ... )  
reviews = @movie.reviews
```

- May be faster:

```
@movie =  
movie.where( ... ).include(:reviews)@movie.reviews.each do |  
review|
```

reviews are already loaded!

Indices

- Speeds up access when searching DB table by column other than primary key
 - e.g. `Movie.where("rating = 'PG'")`
 - Similar to using a hash table
 - alternative is *table scan*—bad!
 - even bigger win if attribute is unique-valued
 - Why not index *every* column?
 - takes up space
 - all indices must be updated when table updated
-

What to index?

- Foreign key columns, eg `movie_id` field in `Reviews` table
- why?
- Columns that appear in `where()` clauses of ActiveRecord queries
- Columns on which you sort
- Use `rails_indexes` gem (on GitHub) to help identify missing indices (and unnecessary ones!)





How much does indexing help?

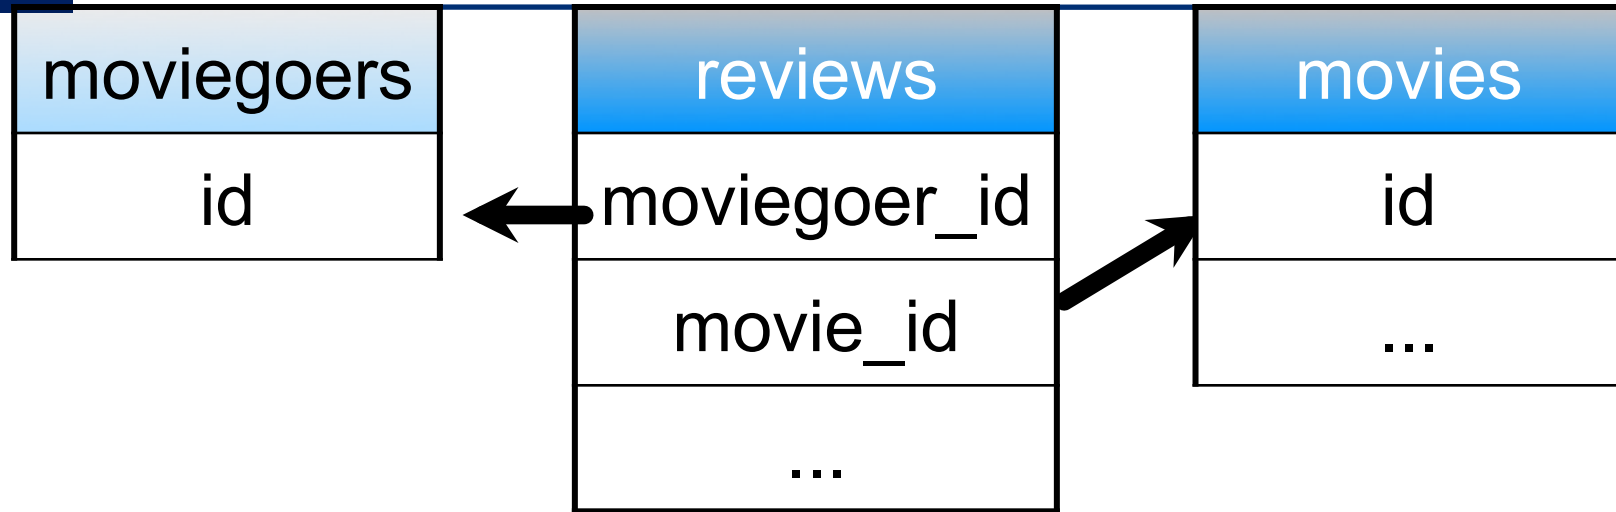
# of reviews:	2000	20,000	200,000
Read 100, no indices	0.94	1.33	5.28
Read 100, FK indices	0.57	0.63	0.65
Performance	166%	212%	808%

200,000	
Create 1K, no indices	9.69
Create 1K, all indices	11.30
Performance	-17%

Suppose Movie has many Moviegoers through Reviews. Which column index would MOST help speed up the query `fans = @movie.moviegoers`

- ❑ `movies.review_id`
- ❑ `reviews.movie_id`
- ❑ `reviews.moviegoer_id`
- ❑ `moviegoers.review_id`

has_many :through



■ **moviegoer:**

has_many :reviews **has_many :movies, :through => :reviews**

■ **movie:** **has_many :reviews**

has_many :moviegoers, :through => :reviews

■ **reviews:** **belongs_to :moviegoer**
belongs_to :movie



Common Attacks on the App

1. Eavesdropping
2. Man-in-the-middle/Session hijack
3. SQL injection
4. Cross-site request forgery (CSRF)
5. Cross-site scripting (XSS)
6. Mass-assignment of sensitive attributes

...more in book



SSL (Secure Sockets Layer)

- Idea: *encrypt* HTTP traffic to foil eavesdroppers
 - Problem: to create a *secure channel*, two parties need to *share a secret* first
 - But on the Web, the two parties don't know each other
 - Solution: *public key cryptography* (Rivest, Shamir, Adelman shared 2002 Turing Award)
-



What SSL Does, and Doesn't

- Each principal has a *key* of 2 matched parts
 - public part: everyone can know it
 - private part: principal keeps secret
 - given one part, cannot deduce the other
 - Key mechanism: *encryption* by one key requires *decryption* by the other
 - If a message can be decrypted with Bob's public key, then Bob must have used created it
 - If I use Bob's public key to create a message, only Bob can read it
-



How SSL works (simplified)

1. Bob.com proves identity to CA
2. CA uses its *private* key to create a “cert” tying this identity to domain name “bob.com”
3. Cert is installed on Bob.com’s server
4. Browser visits <http://bob.com>
5. CA’s public keys *built into browser*, so can check if cert matches hostname
6. *Diffie-Hellman key exchange* is used to bootstrap an encrypted channel for further communication

Use Rails [force_ssl](#) method to force some or all actions to use SSL



What it Does and Doesn't Do

- Assures browser that bob.com is legit
- Prevents eavesdroppers from reading HTTP traffic between browser & bob.com
- Creates additional work for server!

DOES NOT:

- Assure server of who the user is
 - Say anything about what happens to sensitive data *after* it reaches server
 - Say anything about whether server is vulnerable to other server attacks
 - Protect browser from malware if server is evil
-

SQL Injection

- View: = text_field_tag 'name'
- App: Moviegoer.where("name='#{params[:name]}'"")
- Evil user fills in: **BOB'); DROP TABLE moviegoers; --**
- SELECT * FROM moviegoers WHERE (name='BOB');** **DROP TABLE moviegoers; --'**
- Solution: Moviegoer.where("name=?", params[:name])



Cross-site Request Forgery

1. Alice logs into bank.com, now has cookie
2. Alice goes to blog.evil.com
3. Page contains: ``
4. evil.com harvests Alice's personal info

Solutions:

1. (weak) check Referer field in HTTP header
 2. (strong) include *session nonce* with every request
 - `csrf_meta_tags` in layouts/application.html.haml
 - `protect_from_forgery` in ApplicationController
 - Rails form helpers automatically include nonce in forms
-

If a site has a valid SSL certificate from a trusted CA, which of the following is true:

- i) The site is probably not “masquerading” as an impostor of a real site
- ii) CSRF + SQL injection are harder to mount against it
- iii) Your data is secure once it reaches the site

- ☐ (i) only
- ☐ (i) & (ii) only
- ☐ (ii) & (iii) only
- ☐ (i), (ii) & (iii)



Optimizing Prematurely or Without Measurements

- Speed is a feature that users expect
- 99%ile (eg), not “average”
- *Horizontal scaling* >> per-machine performance, *but* lots of ways things can slow down
- Monitoring is your friend: measure twice, cut once
- More: railslab.newrelic.com/scaling-rails



“Mine is a 3-tier app on cloud computing, so it will scale”

- Database is particularly hard to scale
- Even if you do, still want to get “expensive” operations out of the way of your SLO
- One help: cache at many levels
 - whole page, fragment, query
- Cache *expiration* is a crosscutting concern
- RoR support for crosscutting concerns allows you to specify it declaratively
- Use PaaS for as long as you can



“My small site isn’t a target”

- Hackers may be after your users, not your data
 - Like performance, security is a *crosscutting concern*—hard to add after the fact
 - Stay current with best practices and tools—you’re unlikely to do better by rolling your own
 - Prepare for catastrophe: keep regular backups of site and database
-

Your users are sporadically complaining that your site is slow, yet New Relic reports low traffic levels and low CPU utilization. What is the likely cause?

- Not enough Heroku “dynos”, so requests occasionally get “backed up”
- Some queries are unusually slow because you’re sharing DB with other apps
- Some views take unusually long to render in certain browsers (eg, due to JavaScript)
- It could be any of these/Not enough information